

# Architecture Document: Essay Grading System (Fine-Tuned Qwen + LangGraph Framework)

---

## 1. Overview

The Essay Grading System is an **AI-driven pipeline** that evaluates student essays against multiple rubric dimensions. It combines a **fine-tuned Qwen language model** with **custom mapping layers and weighted scoring logic** to provide both **analytic feedback** and a **final aggregated score**.

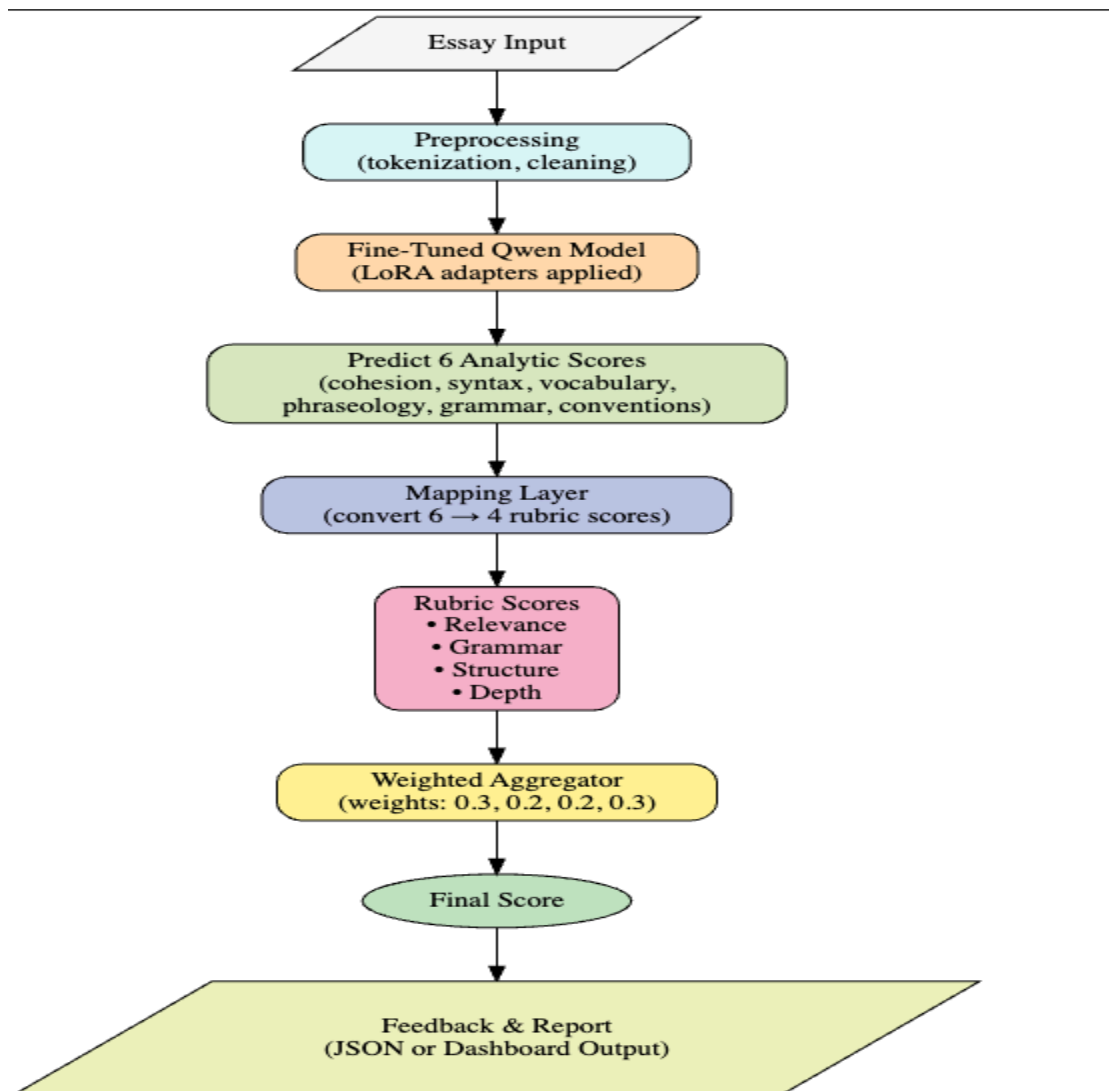
---

## 2. System Objectives

- Automate essay evaluation at scale.
  - Provide **transparent, rubric-aligned scores** instead of a “black-box” grade.
  - Leverage **fine-tuned open-source LLMs (Qwen)** for cost efficiency and flexibility.
  - Allow easy integration with dashboards, APIs, or LMS platforms.
- 

## 3. High-Level Architecture

Flow Diagram



### Steps:

1. **Essay Input**
  - Student submits an essay.
  - Input collected via API, web app, or batch upload.
2. **Preprocessing**
  - Tokenization, cleaning, trimming to model max length.
  - Ensures input fits LLM constraints and removes noise.
3. **Fine-Tuned Qwen Model (LoRA Adapters)**
  - Base: [Qwen/Qwen2.5-3B-Instruct](#).
  - Fine-tuned using **Feedback Prize – ELL dataset** via **QLoRA**.
  - Outputs **six analytic scores**: cohesion, syntax, vocabulary, phraseology, grammar, conventions.
  - Chosen because:

- Open-source (no API costs like GPT-4o).
  - Fine-tuned with domain data → improved grading consistency.
  - LoRA makes it feasible on consumer GPUs (memory-efficient).
4. **Mapping Layer (6 → 4 Rubric Scores)**
- Converts ELL's six dimensions into four rubric categories:
    - **Relevance Score** ← cohesion + vocabulary + phraseology.
    - **Grammar Score** ← grammar (+ conventions weighted).
    - **Structure Score** ← cohesion + syntax.
    - **Depth Score** ← vocabulary + phraseology + syntax.
  - Justification: matches your original framework's design while leveraging analytic outputs.
5. **Rubric Scores Output**
- Each score normalized to [0, 1].
  - Provides interpretable feedback across multiple dimensions.
6. **Weighted Aggregator**
- Combines the four rubric scores into a **final score** using weights:
    - Relevance: **0.3**
    - Grammar: **0.2**
    - Structure: **0.2**
    - Depth: **0.3**
  - Justification: higher weight to **Relevance** and **Depth** as they reflect essay content/insight.
7. **Final Score**
- A single numeric grade (0–1, scaled to 0–100 if needed).
  - Used for comparative ranking or direct grading.
8. **Feedback & Reporting**
- JSON output for programmatic integration.
  - Optional dashboard visualizations:
    - Spider/Radar chart for rubric scores.
    - Histogram of final scores across a class.
  - Enables both **teachers** (analytics) and **students** (feedback).
- 

## 4. Technology Stack

- **Model:** Qwen2.5-3B-Instruct (fine-tuned with LoRA).
  - **Training Framework:** Hugging Face `transformers`, `trl`, `peft`, `bitsandbytes`.
  - **Workflow Framework:** LangGraph for modular, DAG-based grading logic.
  - **Storage/Interface:** JSON output; can be plugged into Flask, FastAPI, or LMS.
  - **Visualization:** Graphviz for architecture diagrams, dashboards for analytics.
-

## 5. Why This Architecture

- **Fine-Tuned Qwen:** Free to Use SLM.
  - **LoRA Adapters:** efficient fine-tuning within Kaggle/Colab GPU limits.
  - **Mapping Layer:** bridges dataset labels → original rubric requirements.
  - **Weighted Aggregation:** preserves grading philosophy, customizable for different institutions.
  - **Modular Flow (LangGraph):** easy to extend (e.g., add plagiarism check, coherence analysis).
- 

## 6. Future Enhancements

- **Multi-task Training:** directly fine-tune Qwen to predict all four rubric scores instead of mapping.
  - **Explainability:** add natural-language feedback per score.
  - **Calibration:** align model outputs with human graders using linear regression or Platt scaling.
  - **Deployment:** package as a REST API or microservice for LMS integration.
-